

1. Business Problem

1.1 Problem Description

Netflix is all about connecting people to the movies they love. To help customers find those movies, they developed world-class movie recommendation system: CinematchSM. Its job is to predict whether someone will enjoy a movie based on how much they liked or disliked other movies. Netflix use those predictions to make personal movie recommendations based on each customer's unique tastes. And while **Cinematch** is doing pretty well, it can always be made better.

Now there are a lot of interesting alternative approaches to how Cinematch works that netflix haven't tried. Some are described in the literature, some aren't. We're curious whether any of these can beat Cinematch by making better predictions. Because, frankly, if there is a much better approach it could make a big difference to our customers and our business.

Credits: <https://www.netflixprize.com/rules.html>

1.2 Problem Statement

Netflix provided a lot of anonymous rating data, and a prediction accuracy bar that is 10% better than what Cinematch can do on the same training data set. (Accuracy is a measurement of how closely predicted ratings of movies match subsequent actual ratings.)

1.3 Sources

- <https://www.netflixprize.com/rules.html>
- <https://www.kaggle.com/netflix-inc/netflix-prize-data>
- Netflix blog: <https://medium.com/netflix-techblog/netflix-recommendations-beyond-the-5-stars-part-1-55838468f429> (very nice blog)
- surprise library: <http://surpriselib.com/> (we use many models from this library)
- surprise library doc: http://surprise.readthedocs.io/en/stable/getting_started.html (we use many models from this library)
- installing surprise: <https://github.com/NicolasHug/Surprise#installation>
- Research paper: <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf> (most of our work was inspired by this paper)
- SVD Decomposition : <https://www.youtube.com/watch?v=P5mlg91as1c>

1.4 Real world/Business Objectives and constraints

Objectives:

1. Predict the rating that a user would give to a movie that he has not yet rated.
2. Minimize the difference between predicted and actual rating (RMSE and MAPE)

Constraints:

1. Some form of interpretability.

2. Machine Learning Problem

2.1 Data

2.1.1 Data Overview

Get the data from : <https://www.kaggle.com/netflix-inc/netflix-prize-data/data>

Data files :

- combined_data_1.txt
- combined_data_2.txt
- combined_data_3.txt
- combined_data_4.txt
- movie_titles.csv

The first line of each file [combined_data_1.txt, combined_data_2.txt, combined_data_3.txt, combined_data_4.txt] contains the movie id followed by a colon. Each subsequent line in the file corresponds to a rating from a customer and its date in the following format:

CustomerID,Rating,Date

MovieIDs range from 1 to 17770 sequentially.

CustomerIDs range from 1 to 2649429, with gaps. There are 480189 users.

Ratings are on a five star (integral) scale from 1 to 5.

Dates have the format YYYY-MM-DD.

2.1.2 Example Data point

```
1:
1488844,3,2005-09-06
822109,5,2005-05-13
885013,4,2005-10-19
30878,4,2005-12-26
823519,3,2004-05-03
893988,3,2005-11-17
124105,4,2004-08-05
1248029,3,2004-04-22
1842128,4,2004-05-09
2238063,3,2005-05-11
1503895,4,2005-05-19
2207774,5,2005-06-06
2590061,3,2004-08-12
2442,3,2004-04-14
543865,4,2004-05-28
1209119,4,2004-03-23
804919,4,2004-06-10
1086807,3,2004-12-28
1711859,4,2005-05-08
372233,5,2005-11-23
1080361,3,2005-03-28
1245640,3,2005-12-19
558634,4,2004-12-14
2165002,4,2004-04-06
1181550,3,2004-02-01
1227322,4,2004-02-06
427928,4,2004-02-26
814701,5,2005-09-29
808731,4,2005-10-31
662870,5,2005-08-24
337541,5,2005-03-23
786312,3,2004-11-16
1133214,4,2004-03-07
1537427,4,2004-03-29
1209954,5,2005-05-09
```

```
2381599,3,2005-09-12
525356,2,2004-07-11
1910569,4,2004-04-12
2263586,4,2004-08-20
2421815,2,2004-02-26
1009622,1,2005-01-19
1481961,2,2005-05-24
401047,4,2005-06-03
2179073,3,2004-08-29
1434636,3,2004-05-01
93986,5,2005-10-06
1308744,5,2005-10-29
2647871,4,2005-12-30
1905581,5,2005-08-16
2508819,3,2004-05-18
1578279,1,2005-05-19
1159695,4,2005-02-15
2588432,3,2005-03-31
2423091,3,2005-09-12
470232,4,2004-04-08
2148699,2,2004-06-05
1342007,3,2004-07-16
466135,4,2004-07-13
2472440,3,2005-08-13
1283744,3,2004-04-17
1927580,4,2004-11-08
716874,5,2005-05-06
4326,4,2005-10-29
```

2.2 Mapping the real world problem to a Machine Learning Problem

2.2.1 Type of Machine Learning Problem

For a given movie and user we need to predict the rating would be given by him/her to the movie.

The given problem is a Recommendation problem

It can also be seen as a Regression problem

2.2.2 Performance metric

- Mean Absolute Percentage Error: https://en.wikipedia.org/wiki/Mean_absolute_percentage_error
- Root Mean Square Error: https://en.wikipedia.org/wiki/Root-mean-square_deviation

2.2.3 Machine Learning Objective and Constraints

1. Minimize RMSE.
2. Try to provide some interpretability.

In [1]:

```
# this is just to know how much time will it take to run this entire ipython notebook
from datetime import datetime
# globalstart = datetime.now()
import pandas as pd
import numpy as np
import matplotlib
matplotlib.use('nbagg')

import matplotlib.pyplot as plt
plt.rcParams.update({'figure.max_open_warning': 0})
```

```

import seaborn as sns
#sns.set_style('whitegrid')
import os
from scipy import sparse
from scipy.sparse import csr_matrix

from sklearn.decomposition import TruncatedSVD
from sklearn.metrics.pairwise import cosine_similarity
import random

```

3. Exploratory Data Analysis

3.1 Preprocessing

3.1.1 Converting / Merging whole data to required format: u_i, m_j, r_ij

In [2]:

```

start = datetime.now()
if not os.path.isfile('data.csv'):
    # Create a file 'data.csv' before reading it
    # Read all the files in netflix and store them in one big file('data.csv')
    # We re reading from each of the four files and appendig each rating to a global file
    'train.csv'
    data = open('data.csv', mode='w')

    row = list()
    files=['combined_data_1.txt','combined_data_2.txt',
           'combined_data_3.txt', 'combined_data_4.txt']
    for file in files:
        print("Reading ratings from {}".format(file))
        with open(file) as f:
            for line in f:
                del row[:] # you don't have to do this.
                line = line.strip()
                if line.endswith(':'):
                    # All below are ratings for this movie, until another movie appears.
                    movie_id = line.replace(':', '')
                else:
                    row = [x for x in line.split(',') ]
                    row.insert(0, movie_id)
                    data.write(','.join(row))
                    data.write('\n')

            print("Done.\n")
    data.close()
print('Time taken :', datetime.now() - start)

```

Time taken : 0:00:00

In [3]:

```

print("creating the dataframe from data.csv file..")
df = pd.read_csv('data.csv', sep=',',
                 names=['movie', 'user', 'rating', 'date'])
df.date = pd.to_datetime(df.date)
print('Done.\n')

# we are arranging the ratings according to time.
print('Sorting the dataframe by date..')
df.sort_values(by='date', inplace=True)
print('Done..')

```

creating the dataframe from data.csv file..
Done.

Sorting the dataframe by date..
Done..

In [4]:

```
df.head()
```

Out[4]:

	movie	user	rating	date
56431994	10341	510180	4	1999-11-11
9056171	1798	510180	5	1999-11-11
58698779	10774	510180	3	1999-11-11
48101611	8651	510180	2	1999-11-11
81893208	14660	510180	2	1999-11-11

In [5]:

```
df.describe() ['rating']
```

Out[5]:

```
count      1.004805e+08
mean        3.604290e+00
std          1.085219e+00
min          1.000000e+00
25%          3.000000e+00
50%          4.000000e+00
75%          4.000000e+00
max          5.000000e+00
Name: rating, dtype: float64
```

3.1.2 Checking for NaN values

In [6]:

```
# just to make sure that all Nan containing rows are deleted..
print("No of Nan values in our dataframe : ", sum(df.isnull().any()))
```

No of Nan values in our dataframe : 0

3.1.3 Removing Duplicates

In [6]:

```
dup_bool = df.duplicated(['movie','user','rating'])
dups = sum(dup_bool) # by considering all columns..( including timestamp)
print("There are {} duplicate rating entries in the data..".format(dups))
```

There are 0 duplicate rating entries in the data..

3.1.4 Basic Statistics (#Ratings, #Users, and #Movies)

In [7]:

```
print("Total data ")
print("-"*50)
print("\nTotal no of ratings :",df.shape[0])
print("Total No of Users      :", len(np.unique(df.user)))
```

```
print("Total No of movies :", len(np.unique(df.movie)))
```

Total data

```
-----  
Total no of ratings : 100480507  
Total No of Users   : 480189  
Total No of movies  : 17770
```

In [8]:

```
print(df.shape)
```

```
(100480507, 4)
```

3.2 Splitting data into Train and Test(80:20)

In [2]:

```
if not os.path.isfile('train.csv'):  
    # create the dataframe and store it in the disk for offline purposes..  
    df.iloc[:int(df.shape[0]*0.80)].to_csv("train.csv", index=False)  
  
if not os.path.isfile('test.csv'):  
    # create the dataframe and store it in the disk for offline purposes..  
    df.iloc[int(df.shape[0]*0.80):].to_csv("test.csv", index=False)  
  
train_df = pd.read_csv("train.csv", parse_dates=['date'])  
test_df = pd.read_csv("test.csv")
```

3.2.1 Basic Statistics in Train data (#Ratings, #Users, and #Movies)

In [10]:

```
# movies = train_df.movie.value_counts()  
# users = train_df.user.value_counts()  
print("Training data ")  
print("-"*50)  
print("\nTotal no of ratings :", train_df.shape[0])  
print("Total No of Users   :", len(np.unique(train_df.user)))  
print("Total No of movies  :", len(np.unique(train_df.movie)))
```

Training data

```
-----  
Total no of ratings : 67725000  
Total No of Users   : 356718  
Total No of movies  : 16902
```

3.2.2 Basic Statistics in Test data (#Ratings, #Users, and #Movies)

In [11]:

```
print("Test data ")  
print("-"*50)  
print("\nTotal no of ratings :", test_df.shape[0])  
print("Total No of Users   :", len(np.unique(test_df.user)))  
print("Total No of movies  :", len(np.unique(test_df.movie)))
```

Test data

```
-----  
Total no of ratings : 20096102  
Total No of Users   : 349312  
Total No of movies  : 17757
```

In [12]:

```
print(train_df.head)
```

```
<bound method NDFrame.head of
0      10341  510180    4 1999-11-11    user  rating    date
1       1798  510180    5 1999-11-11
2      10774  510180    3 1999-11-11
3       8651  510180    2 1999-11-11
4      14660  510180    2 1999-11-11
...      ...      ...    ...    ...
67724995 12828 1168143    3 2005-05-23
67724996 13392 2259724    4 2005-05-23
67724997  1470 1233482    3 2005-05-23
67724998   152 2009801    3 2005-05-23
67724999   8990 556844    3 2005-05-23
```

```
[67725000 rows x 4 columns]>
```

3.3 Exploratory Data Analysis on Train data

In [3]:


```
# method to make y-axis more readable
def human(num, units = 'M'):
    units = units.lower()
    num = float(num)
    if units == 'k':
        return str(num/10**3) + " K"
    elif units == 'm':
        return str(num/10**6) + " M"
    elif units == 'b':
        return str(num/10**9) + " B"
```


3.3.1 Distribution of ratings

In [4]:

```
fig, ax = plt.subplots()
plt.title('Distribution of ratings over Training dataset', fontsize=15)
sns.countplot(train_df.rating)
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
ax.set_ylabel('No. of Ratings(Millions)')

plt.show()
```





In []:

```
for item in ax.get_yticks():
    print(item)
```

Add new column (week day) to the data set for analysis.

In [5]:

```
# It is used to skip the warning "'SettingWithCopyWarning'.."
pd.options.mode.chained_assignment = None # default='warn'

train_df['day_of_week'] = train_df.date.dt.day_name

train_df.tail()
```


Out[5]:


	movie	user	rating	date	day_of_week
67724995	12828	1168143	3	2005-05-23	<bound method PandasDelegate._add_delegate_acc...
67724996	13392	2259724	4	2005-05-23	<bound method PandasDelegate._add_delegate_acc...
67724997	1470	1233482	3	2005-05-23	<bound method PandasDelegate._add_delegate_acc...
67724998	152	2009801	3	2005-05-23	<bound method PandasDelegate._add_delegate_acc...
67724999	8990	556844	3	2005-05-23	<bound method PandasDelegate._add_delegate_acc...

3.3.2 Number of Ratings per a month

In [6]:

```
ax = train_df.resample('m', on='date')['rating'].count().plot()
ax.set_title('No of ratings per month (Training data)')
plt.xlabel('Month')
plt.ylabel('No of ratings(per month)')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```





3.3.3 Analysis on the Ratings given by user

In [7]:

```
no_of Rated movies per user = train_df.groupby(by='user') ['rating'].count().sort_values(ascending=False)

no_of Rated movies per user.head()
```

Out[7]:

```
user
305344      16509
2439493     15203
387418      14512
1639792      9767
1461435      8832
Name: rating, dtype: int64
```

In [8]:

```
fig = plt.figure(figsize=plt.figaspect(.5))
#import statsmodels.api as sm
ax1 = plt.subplot(121)
#dens = sm.nonparametric.KDEUnivariate(no_of Rated movies per user,shade=True, ax=ax1)
#dens.fit()
sns.kdeplot(no_of Rated movies per user, shade=True, ax=ax1)
#plt.plot(dens.pdf)
plt.xlabel('No of ratings by user')
plt.title("PDF")

ax2 = plt.subplot(122)

counts, bin_edges = np.histogram(no_of Rated movies per user, bins=20, normed=True)


# Now find the cdf
cdf = np.cumsum(counts)


ax2.plot(bin_edges[1:],cdf)
#ax2.fill_between(bin_edges[:-1], cdf, 0,facecolor="lightblue",alpha=0.5)
plt.xlabel('No of ratings by user')
plt.legend('rating')
plt.title('CDF')

plt.show()
```

c:\users\admin\anaconda3\lib\site-packages\ipykernel_launcher.py:13: VisibleDeprecationWarning: Passing 'normed=True' on non-uniform bins has always been broken, and computes neither the probability density function nor the probability mass function. The result is only correct if the bins are uniform, when density=True will produce the same result anyway. The argument will be removed in a future version of numpy.

```
del sys.path[0]
```





In [9]:

```
no_of Rated movies per user.describe()
```

Out[9]:

```
count    356718.000000
mean      189.855853
std       283.837049
min        1.000000
25%       27.000000
50%       84.000000
75%      236.000000
max     16509.000000
Name: rating, dtype: float64
```

There, is something interesting going on with the quantiles..

In [10]:


```
quantiles = no_of Rated movies per user.quantile(np.arange(0,1.01,0.01), interpolation='higher')
```

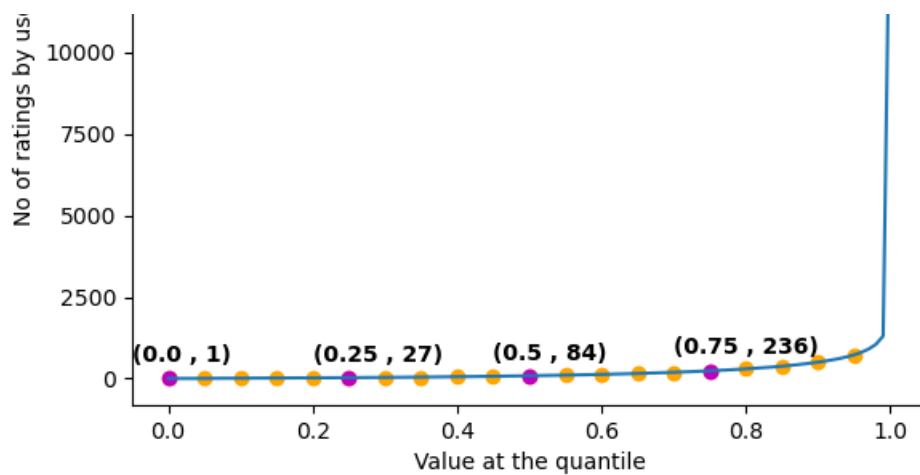
In [11]:

```
plt.title("Quantiles and their Values")
quantiles.plot()
# quantiles with 0.05 difference
plt.scatter(x=quantiles.index[::5], y=quantiles.values[::5], c='orange', label="quantiles with 0.05 intervals")
# quantiles with 0.25 difference
plt.scatter(x=quantiles.index[::25], y=quantiles.values[::25], c='m', label = "quantiles with 0.25 intervals")
plt.ylabel('No of ratings by user')
plt.xlabel('Value at the quantile')
plt.legend(loc='best')

# annotate the 25th, 50th, 75th and 100th percentile values....
for x,y in zip(quantiles.index[::25], quantiles.values[::25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500)
                , fontweight='bold')

plt.show()
```





In [12]:

```
quantiles[::5]
```

Out[12]:

```
0.00      1
0.05      4
0.10     10
0.15     16
0.20     21
0.25     27
0.30     35
0.35     44
0.40     55
0.45     69
0.50     84
0.55    104
0.60    127
0.65    155
0.70    191
0.75    236
0.80    296
0.85    378
0.90    503
0.95    728
1.00   16509
Name: rating, dtype: int64
```

how many ratings at the last 5% of all ratings??

In [13]:

```
print('\n No of ratings at last 5 percentile : {}'.format(sum(no_of Rated movies per user >= 749)
) )
```

No of ratings at last 5 percentile : 16830

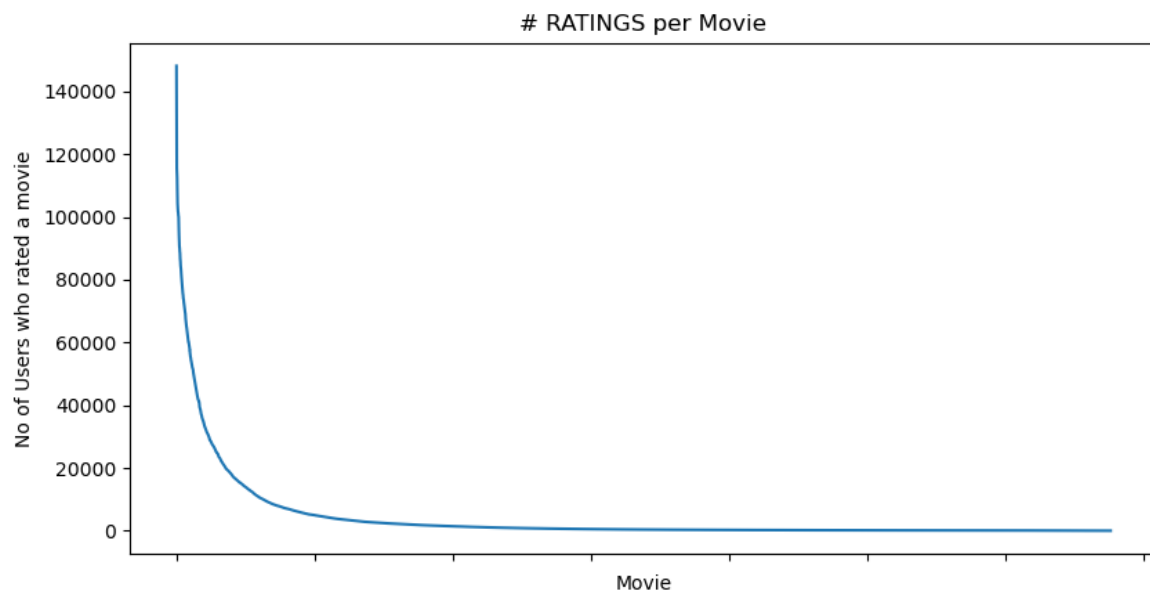
3.3.4 Analysis of ratings of a movie given by a user

In [14]:

```
no_of_ratings_per_movie = train_df.groupby(by='movie')
['rating'].count().sort_values(ascending=False)

fig = plt.figure(figsize=plt.figaspect(.5))
ax = plt.gca()
plt.plot(no_of_ratings_per_movie.values)
plt.title('# RATINGS per Movie')
plt.xlabel('Movie')
plt.ylabel('No of Users who rated a movie')
ax.set_xticklabels([])
```

```
ax.set_xticklabels([''],
plt.show()
```

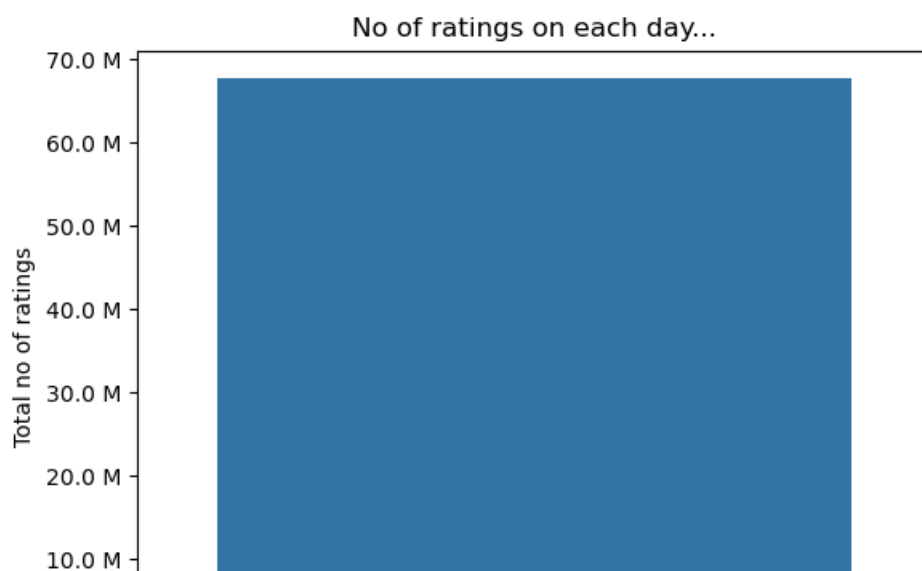


- It is very skewed.. just like number of ratings given per user.
 - There are some movies (which are very popular) which are rated by huge number of users.
 - But most of the movies (like 90%) got some hundreds of ratings.

3.3.5 Number of ratings on each day of the week

In [15]:

```
fig, ax = plt.subplots()
sns.countplot(x='day_of_week', data=train_df, ax=ax)
plt.title('No of ratings on each day...')
plt.ylabel('Total no of ratings')
plt.xlabel('')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```



0.0 M

essors.<locals>._create_delegator_method.<locals>.f of <pandas.core.indexes.accessors.C

In [20]:

```
start = datetime.now()
fig = plt.figure(figsize=plt.figaspect(.45))
sns.boxplot(y='rating', x='day_of_week', data=train_df)
plt.show()
print(datetime.now() - start)
```

0:00:19.505394

In [21]:

```
avg_week_df = train_df.groupby(by=['day_of_week'])['rating'].mean()
print(" AVerage ratings")
print("-"*30)
print(avg_week_df)
print("\n")
```

AVerage ratings

day_of_week

<bound method PandasDelegate._add_delegate_accessors.<locals>._create_delegator_method.<locals>.f
of <pandas.core.indexes.accessors.DatetimeProperties object at 0x0000013F42C115C8>> 3.563122
Name: rating, dtype: float64

3.3.6 Creating sparse matrix from data frame

3.3.6.1 Creating sparse matrix from train data frame

In [18]:

```
start = datetime.now()
if os.path.isfile('train_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    train_sparse_matrix = sparse.load_npz('train_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    train_sparse_matrix = sparse.csr_matrix((train_df.rating.values, (train_df.user.values,
                                                                    train_df.movie.values)),)

    print('Done. It\'s shape is : (user, movie) : ', train_sparse_matrix.shape)
    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz("train_sparse_matrix.npz", train_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....

DONE..

0:00:04.052068

The Sparsity of Train Sparse Matrix

In [19]:

```
us,mv = train_sparse_matrix.shape
elem = train_sparse_matrix.count_nonzero()

print("Sparsity Of Train matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )
```

Sparsity Of Train matrix : 99.85615833640739 %

3.3.6.2 Creating sparse matrix from test data frame

In [22]:

```
start = datetime.now()
if os.path.isfile('test_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    test_sparse_matrix = sparse.load_npz('test_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    test_sparse_matrix = sparse.csr_matrix((test_df.rating.values, (test_df.user.values,
                                                                    test_df.movie.values)))

    print('Done. It\'s shape is : (user, movie) : ',test_sparse_matrix.shape)
    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz("test_sparse_matrix.npz", test_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....
DONE..
0:00:01.479840

The Sparsity of Test data Matrix

In [23]:

```
us,mv = test_sparse_matrix.shape
elem = test_sparse_matrix.count_nonzero()

print("Sparsity Of Test matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )
```

Sparsity Of Test matrix : 99.95731772988694 %

3.3.7 Finding Global average of all movie ratings, Average rating per user, and Average rating per movie

In [24]:

```
# get the user averages in dictionary (key: user_id/movie_id, value: avg rating)

def get_average_ratings(sparse_matrix, of_users):

    # average ratings of user/axes
    ax = 1 if of_users else 0 # 1 - User axes, 0 - Movie axes

    # ".A1" is for converting Column_Matrix to 1-D numpy array
    sum of ratings = sparse_matrix.sum(axis=ax).A1
```

```

# Boolean matrix of ratings ( whether a user rated that movie or not)
is_rated = sparse_matrix!=0

# no of ratings that each user OR movie..
no_of_ratings = is_rated.sum(axis=ax).A1

# max_user and max_movie ids in sparse matrix
u,m = sparse_matrix.shape
# create a dictionary of users and their average ratings..
average_ratings = { i : sum_of_ratings[i]/no_of_ratings[i]
                    for i in range(u if of_users else m)
                    if no_of_ratings[i] !=0}

# return that dictionary of average ratings
return average_ratings

```

3.3.7.1 finding global average of all movie ratings

In [25]:

```

train_averages = dict()
# get the global average of ratings in our train set.
train_global_average = train_sparse_matrix.sum()/train_sparse_matrix.count_nonzero()
train_averages['global'] = train_global_average
train_averages

```

Out[25]:

```
{'global': 3.5631217128091546}
```

3.3.7.2 finding average rating per user

In [26]:

```

train_averages['user'] = get_average_ratings(train_sparse_matrix, of_users=True)
print('\nAverage rating of user 10 : ',train_averages['user'][10])

```

Average rating of user 10 : 3.3684210526315788

3.3.7.3 finding average rating per movie

In [27]:

```

train_averages['movie'] = get_average_ratings(train_sparse_matrix, of_users=False)
print('\nAverage rating of movie 15 : ',train_averages['movie'][15])

```

Average rating of movie 15 : 3.2911392405063293

3.3.7.4 PDF's & CDF's of Avg.Ratings of Users & Movies (In Train Data)

In [28]:

```

start = datetime.now()
# draw pdfs for average rating per user and average
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))
fig.suptitle('Avg Ratings per User and per Movie', fontsize=15)

ax1.set_title('Users-Avg-Ratings')
# get the list of average user ratings from the averages dictionary..
user_averages = [rat for rat in train_averages['user'].values()]
counts, bin_edges = np.histogram(user_averages, bins=20, normed=True)

# Now find the cdf
cdf = np.cumsum(counts/sum(counts))
ax1.plot(bin_edges[1:],cdf,label="CDF")

```

```

sns.distplot(user_averages, ax=ax1, hist=False, label="PDF")

ax2.set_title('Movies-Avg-Rating')
# get the list of movie_average_ratings from the dictionary..
movie_averages = [rat for rat in train_averages['movie'].values()]

counts, bin_edges = np.histogram(movie_averages, bins=20, normed=True)

# Now find the cdf
cdf = np.cumsum(counts/sum(counts))

ax2.plot(bin_edges[1:], cdf, label="CDF")
#ax2.fill_between(bin_edges[:-1], cdf, 0, facecolor="lightblue", alpha=0.5)

sns.distplot(movie_averages, ax=ax2, hist=False, label='Pdf')

plt.show()
print(datetime.now() - start)

```

c:\users\admin\anaconda3\lib\site-packages\ipykernel_launcher.py:9: VisibleDeprecationWarning: Passing `normed=True` on non-uniform bins has always been broken, and computes neither the probability density function nor the probability mass function. The result is only correct if the bins are uniform, when density=True will produce the same result anyway. The argument will be removed in a future version of numpy.

c:\users\admin\anaconda3\lib\site-packages\ipykernel_launcher.py:22: VisibleDeprecationWarning: Passing `normed=True` on non-uniform bins has always been broken, and computes neither the probability density function nor the probability mass function. The result is only correct if the bins are uniform, when density=True will produce the same result anyway. The argument will be removed in a future version of numpy.

0:00:00.320214

3.3.8 Cold Start problem

3.3.8.1 Cold Start problem with Users

In [29]:

```

total_users = len(np.unique(df.user))
users_train = len(train_averages['user'])
new_users = total_users - users_train

print('\nTotal number of Users  :', total_users)
print('\nNumber of Users in Train data :', users_train)
print("\nNo of Users that didn't appear in train data: {} ({} %) \n ".format(new_users,

np.round((new_users/total_users)*100, 2)))

```

NameError Traceback (most recent call last)

```

<ipython-input-29-c1633af1f47f> in <module>
----> 1 total_users = len(np.unique(df.user))
      2 users_train = len(train_averages['user'])
      3 new_users = total_users - users_train
      4
      5 print('\nTotal number of Users  :', total_users)

```

NameError: name 'df' is not defined

We might have to handle **new users** (**123471**) who didn't appear in train data.

3.3.8.2 Cold Start problem with Movies

In []:

```
total_movies = len(np.unique(df.movie))
movies_train = len(train_averages['movie'])
new_movies = total_movies - movies_train

print('\nTotal number of Movies :', total_movies)
print('\nNumber of Users in Train data :', movies_train)
print("\nNo of Movies that didn't appear in train data: {} ({} %) \n ".format(new_movies,
np.round((new_movies/total_movies)*100, 2)))
```

We might have to handle **868 movies** (small comparatively) in test data

3.4 Computing Similarity matrices

3.4.1 Computing User-User Similarity matrix

1. Calculating User User Similarity_Matrix is **not very easy**(unless you have huge Computing Power and lots of time) because of number of. users being large.
 - You can try if you want to. Your system could crash or the program stops with **Memory Error**

3.4.1.1 Trying with all dimensions (17k dimensions per user)

In []:

```
from sklearn.metrics.pairwise import cosine_similarity

def compute_user_similarity(sparse_matrix, compute_for_few=False, top = 100, verbose=False, verb_for_n_rows = 20,
                           draw_time_taken=True):
    no_of_users, _ = sparse_matrix.shape
    # get the indices of non zero rows(users) from our sparse matrix
    row_ind, col_ind = sparse_matrix.nonzero()
    row_ind = sorted(set(row_ind)) # we don't have to
    time_taken = list() # time taken for finding similar users for an user..

    # we create rows, cols, and data lists.., which can be used to create sparse matrices
    rows, cols, data = list(), list(), list()
    if verbose: print("Computing top",top,"similarities for each user..")

    start = datetime.now()
    temp = 0

    for row in row_ind[:top] if compute_for_few else row_ind:
        temp = temp+1
        prev = datetime.now()

        # get the similarity row for this user with all other users
        sim = cosine_similarity(sparse_matrix.getrow(row), sparse_matrix).ravel()
        # We will get only the top 'top' most similar users and ignore rest of them..
        top_sim_ind = sim.argsort()[-top:]
        top_sim_val = sim[top_sim_ind]

        # add them to our rows, cols and data
        rows.extend([row]*top)
        cols.extend(top_sim_ind)
        data.extend(top_sim_val)
        time_taken.append(datetime.now().timestamp() - prev.timestamp())
    if verbose:
```

```

        if temp%verb_for_n_rows == 0:
            print("computing done for {} users [ time elapsed : {} ]".format(temp, datetime.now()-start))

# lets create sparse matrix out of these and return it
if verbose: print('Creating Sparse matrix from the computed similarities')
#return rows, cols, data

if draw_time_taken:
    plt.plot(time_taken, label = 'time taken for each user')
    plt.plot(np.cumsum(time_taken), label='Total time')
    plt.legend(loc='best')
    plt.xlabel('User')
    plt.ylabel('Time (seconds)')
    plt.show()

return sparse.csr_matrix((data, (rows, cols)), shape=(no_of_users, no_of_users)), time_taken

```

In []:

```

start = datetime.now()
u_u_sim_sparse, _ = compute_user_similarity(train_sparse_matrix, compute_for_few=True, top = 100,
                                           verbose=True)

print("-"*100)
print("Time taken :",datetime.now()-start)

```

3.4.1.2 Trying with reduced dimensions (Using TruncatedSVD for dimensionality reduction of user vector)

- We have **405,041 users** in our training set and computing similarities between them..(**17K dimensional vector**..) is time consuming..
- From above plot, It took roughly **8.88 sec** for computing similar users for **one user**
- We have **405,041 users** with us in training set.
- $405041 \times 8.88 = 3596764.08\text{sec} = 59946.068\text{ min}$
 - Even if we run on 4 cores parallelly (a typical system now a days), It will still take almost **10 and 1/2 days**.

IDEA: Instead, we will try to reduce the dimensions using SVD, so that **it might speed up the process**...

In []:

```

from datetime import datetime
from sklearn.decomposition import TruncatedSVD

start = datetime.now()

# initialize the algorithm with some parameters..
# All of them are default except n_components. n_itr is for Randomized SVD solver.
netflix_svd = TruncatedSVD(n_components=500, algorithm='randomized', random_state=4)
trunc_svd = netflix_svd.fit_transform(train_sparse_matrix)

print(datetime.now()-start)

```

Here,

- $\sum \rightarrow (\text{netflix_svd.singular_values_})$
- $\bigvee^T \rightarrow (\text{netflix_svd.components_})$
- \bigcup is not returned. instead **Projection_of_X** onto the new vectorspace is returned.
- It uses **randomized svd** internally, which returns **All 3 of them separately**. Use that instead..

In []:

```

expl_var = np.cumsum(netflix_svd.explained_variance_ratio_)

```

```
expl_var = np.cumsum(netflix_svd.explained_variance_ratio_)
```

```
In [ ]:
```

```
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))

ax1.set_ylabel("Variance Explained", fontsize=15)
ax1.set_xlabel("# Latent Facors", fontsize=15)
ax1.plot(expl_var)
# annotate some (latentfactors, expl_var) to make it clear
ind = [1, 2, 4, 8, 20, 60, 100, 200, 300, 400, 500]
ax1.scatter(x = [i-1 for i in ind], y = expl_var[[i-1 for i in ind]], c='#ff3300')
for i in ind:
    ax1.annotate(s="({}, {})".format(i, np.round(expl_var[i-1], 2)), xy=(i-1, expl_var[i-1]),
                xytext = (i+20, expl_var[i-1] - 0.01), fontweight='bold')

change_in_expl_var = [expl_var[i+1] - expl_var[i] for i in range(len(expl_var)-1)]
ax2.plot(change_in_expl_var)

ax2.set_ylabel("Gain in Var_Expl with One Additional LF", fontsize=10)
ax2.yaxis.set_label_position("right")
ax2.set_xlabel("# Latent Facors", fontsize=20)

plt.show()
```

```
In [ ]:
```

```
for i in ind:
    print("({}, {})".format(i, np.round(expl_var[i-1], 2)))
```

I think 500 dimensions is good enough

- By just taking **(20 to 30)** latent factors, explained variance that we could get is **20 %**.
- To take it to **60%**, we have to take **almost 400 latent factors**. It is not fare.
- It basically is the **gain of variance explained**, if we **add one additional latent factor to it**.
- By adding one by one latent factore too it, the **_gain in expained variance** with that addition is decreasing. (Obviously, because they are sorted that way).
- **LHS Graph:**
 - **x** --- (No of latent factos),
 - **y** --- (The variance explained by taking x latent factors)
- **More decrease in the line (RHS graph) :**
 - We are getting more expained variance than before.
- **Less decrease in that line (RHS graph) :**
 - We are not getting benifitted from adding latent factor furthur. This is what is shown in the plots.
- **RHS Graph:**
 - **x** --- (No of latent factors),
 - **y** --- (Gain n Expl_Var by taking one additional latent factor)

```
In [ ]:
```

```
# Let's project our Original U_M matrix into into 500 Dimensional space...
start = datetime.now()
trunc_matrix = train_sparse_matrix.dot(netflix_svd.components_.T)
print(datetime.now()- start)
```

```
In [ ]:
```

```
type(trunc_matrix), trunc_matrix.shape
```

- Let's convert this to actual sparse matrix and store it for future purposes

In []:

```
if not os.path.isfile('trunc_sparse_matrix.npz'):
    # create that sparse matrix
    trunc_sparse_matrix = sparse.csr_matrix(trunc_matrix)
    # Save this truncated sparse matrix for later usage..
    sparse.save_npz('trunc_sparse_matrix', trunc_sparse_matrix)
else:
    trunc_sparse_matrix = sparse.load_npz('trunc_sparse_matrix.npz')
```

In []:

```
trunc_sparse_matrix.shape
```

In []:

```
start = datetime.now()
trunc_u_u_sim_matrix, _ = compute_user_similarity(trunc_sparse_matrix, compute_for_few=True, top=50,
, verbose=True,
                                                    verb_for_n_rows=10)
print("-"*50)
print("time:", datetime.now()-start)
```

: This is taking more time for each user than Original one.

- from above plot, It took almost **12.18** for computing similar users for **one user**
- We have **405041 users** with us in training set.
- { 405041 \times 12.18 ==== 4933399.38 \sec } ==== 82223.323 \min ==== 1370.388716667 \text{ hours} ==== 57.099529861 \text{ days}...
- Even we run on 4 cores parallelly (a typical system now a days), It will still take almost **(14 - 15)** days.
- **Why did this happen...??**
 - Just think about it. It's not that difficult.

----- (sparse & dense.....get it ??)-----

Is there any other way to compute user user similarity..??

-An alternative is to compute similar users for a particular user, whenever required (ie., **Run time**)

- We maintain a binary Vector for users, which tells us whether we already computed or not..
- *****If not*** :**
 - Compute top (let's just say, 1000) most similar users for this given user, and add this to our datastructure, so that we can just access it(similar users) without recomputing it again.
- *****If It is already Computed***:**
 - Just get it directly from our datastructure, which has that information.
 - In production time, We might have to recompute similarities, if it is computed a long time ago. Because user preferences changes over time. If we could maintain some kind of Timer, which when expires, we have to update it (recompute it).
- *****Which datastructure to use*****
 - It is purely implementation dependant.
 - One simple method is to maintain a ****Dictionary Of Dictionaries****.
 - ****key : ** _userid_**
 - **__value__ : _Again a dictionary_**
 - **__key__ : _Similar User_**

- __value__: _Similarity Value_

3.4.2 Computing Movie-Movie Similarity matrix

In [30]:

```
start = datetime.now()
if not os.path.isfile('m_m_sim_sparse.npz'):
    print("It seems you don't have that file. Computing movie_movie similarity...")
    start = datetime.now()
    m_m_sim_sparse = cosine_similarity(X=train_sparse_matrix.T, dense_output=False)
    print("Done..")
    # store this sparse matrix in disk before using it. For future purposes.
    print("Saving it to disk without the need of re-computing it again.. ")
    sparse.save_npz("m_m_sim_sparse.npz", m_m_sim_sparse)
    print("Done..")
else:
    print("It is there, We will get it.")
    m_m_sim_sparse = sparse.load_npz("m_m_sim_sparse.npz")
    print("Done ...")

print("It's a ", m_m_sim_sparse.shape, " dimensional matrix")

print(datetime.now() - start)
```

```
It is there, We will get it.
Done ...
It's a (17771, 17771) dimensional matrix
0:00:33.989069
```

In [31]:

```
m_m_sim_sparse.shape
```

Out[31]:

```
(17771, 17771)
```

- Even though we have similarity measure of each movie, with all other movies, We generally don't care much about least similar movies.
- Most of the times, only top_xxx similar items matters. It may be 10 or 100.
- We take only those top similar movie ratings and store them in a saperate dictionary.

In [32]:

```
movie_ids = np.unique(m_m_sim_sparse.nonzero()[1])
```

In [33]:

```
start = datetime.now()
similar_movies = dict()
for movie in movie_ids:
    # get the top similar movies and store them in the dictionary
    sim_movies = m_m_sim_sparse[movie].toarray().ravel().argsort()[::-1][1:]
    similar_movies[movie] = sim_movies[:100]
print(datetime.now() - start)

# just testing similar movies for movie_15
similar_movies[15]
```

```
0:01:04.165999
```

Out[33]:

```
array([ 8013, 12049, 16528,  8279,  3755, 13105,  5927,  4424, 15144,
        10193,   590, 15571,  9071, 15054,  3973,  9584,  4549,  6349,
         6116, 16402, 14059,  5370, 16455,  9376,  4624,   149, 3022])
```

```

    8110, 10402, 14009, 3370, 10433, 3370, 4024, 149, 3022,
10597, 14920, 6292, 1596, 13980, 9566, 7428, 9460, 11653,
17139, 1253, 17183, 17610, 10011, 11730, 1720, 17590, 4706,
8003, 2187, 9802, 9166, 5865, 598, 5697, 10199, 1942,
15390, 9688, 4513, 11981, 17584, 376, 8988, 10788, 17115,
565, 1690, 13013, 9427, 2818, 16309, 16334, 1510, 4649,
3338, 13931, 7845, 8873, 12762, 17285, 15360, 3706, 8875,
6410, 9558, 7481, 11947, 12954, 9840, 2716, 11175, 7859,
9488, 9969, 11867, 4467, 10319, 5871, 847, 2879, 5921,
14696], dtype=int64)

```

3.4.3 Finding most similar movies using similarity matrix

Does Similarity really works as the way we expected...?

Let's pick some random movie and check for its similar movies....

In [34]:

```

# First Let's load the movie details into soe dataframe..
# movie details are in 'netflix/movie_titles.csv'

movie_titles = pd.read_csv("movie_titles.csv", sep=',', header = None,
                           names=['movie_id', 'year_of_release', 'title'], verbose=True,
                           index_col = 'movie_id', encoding = "ISO-8859-1")

movie_titles.head()

```

Tokenization took: 10.99 ms

Type conversion took: 27.03 ms

Parser memory cleanup took: 0.00 ms

Out[34]:

	year_of_release	title
movie_id		
1	2003.0	Dinosaur Planet
2	2004.0	Isle of Man TT 2004 Review
3	1997.0	Character
4	1994.0	Paula Abdul's Get Up & Dance
5	2004.0	The Rise and Fall of ECW

Similar Movies for 'Vampire Journals'

In [35]:

```

mv_id = 67

print("\nMovie ---->",movie_titles.loc[mv_id].values[1])

print("\nIt has {} Ratings from users.".format(train_sparse_matrix[:,mv_id].getnnz()))

print("\nWe have {} movies which are similarto this  and we will get only top most..".format(m_m_sim_sparse[:,mv_id].getnnz()))

```

Movie ----> Vampire Journals

It has 250 Ratings from users.

We have 16729 movies which are similarto this and we will get only top most..

In [36]:

```

similarities = m_m_sim_sparse[mv_id].toarray().ravel()

```

```
similar_indices = similarities.argsort()[::-1][1:]

similarities[similar_indices]

sim_indices = similarities.argsort()[::-1][1:] # It will sort and reverse the array and ignore its
similarity (ie.,1)
# and return its indices(movie_ids)
```

In [37]:

```
plt.plot(similarities[sim_indices], label='All the ratings')
plt.plot(similarities[sim_indices[:100]], label='top 100 similar movies')
plt.title("Similar Movies of {} (movie_id)".format(mv_id), fontsize=20)
plt.xlabel("Movies (Not Movie_Ids)", fontsize=15)
plt.ylabel("Cosine Similarity", fontsize=15)
plt.legend()
plt.show()
```

Top 10 similar movies

In [38]:

```
movie_titles.loc[sim_indices[:10]]
```

Out [38]:

movie_id	year_of_release	title
323	1999.0	Modern Vampires
4044	1998.0	Subspecies 4: Bloodstorm
1688	1993.0	To Sleep With a Vampire
12053	1993.0	Dracula Rising
1900	1997.0	Club Vampire
13962	2001.0	Dracula: The Dark Prince
4667	1996.0	Vampirella
16279	2002.0	Vampires: Los Muertos
13873	2001.0	The Breed
7573	2000.0	Kindred: The Embraced: The Complete Vampire Co...

Similarly, we can **find similar users** and compare how similar they are.

4. Machine Learning Models

In [39]:

```
def get_sample_sparse_matrix(sparse_matrix, no_users, no_movies, path, verbose = True):
    """
    It will get it from the 'path' if it is present or It will create
    and store the sampled sparse matrix in the path specified.
    """
```

```

# get (row, col) and (rating) tuple from sparse_matrix...
row_ind, col_ind, ratings = sparse.find(sparse_matrix)
users = np.unique(row_ind)
movies = np.unique(col_ind)

print("Original Matrix : (users, movies) -- ({ } { })".format(len(users), len(movies)))
print("Original Matrix : Ratings -- { }\n".format(len(ratings)))

# It just to make sure to get same sample everytime we run this program..
# and pick without replacement....
np.random.seed(15)
sample_users = np.random.choice(users, no_users, replace=False)
sample_movies = np.random.choice(movies, no_movies, replace=False)
# get the boolean mask or these sampled_items in originl row/col_inds..
mask = np.logical_and( np.isin(row_ind, sample_users),
                        np.isin(col_ind, sample_movies) )

sample_sparse_matrix = sparse.csr_matrix((ratings[mask], (row_ind[mask], col_ind[mask])),
                                         shape=(max(sample_users)+1, max(sample_movies)+1))

if verbose:
    print("Sampled Matrix : (users, movies) -- ({ } { })".format(len(sample_users), len(sample_movies)))
    print("Sampled Matrix : Ratings --", format(ratings[mask].shape[0]))

print('Saving it into disk for furthur usage..')
# save it into disk
sparse.save_npz(path, sample_sparse_matrix)
if verbose:
    print('Done..\n')

return sample_sparse_matrix

```

4.1 Sampling Data

4.1.1 Build sample train data from the train data

In [40]:

```

start = datetime.now()
path = "sample_train_sparse_matrix_new.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    sample_train_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 10k users and 1k movies from available data
    sample_train_sparse_matrix = get_sample_sparse_matrix(train_sparse_matrix, no_users=15000, no_movies=1200,
                                                         path = path)

print(datetime.now() - start)

```

```

Original Matrix : (users, movies) -- (356718 16902)
Original Matrix : Ratings -- 67725000

```

```

Sampled Matrix : (users, movies) -- (15000 1200)
Sampled Matrix : Ratings -- 199566
Saving it into disk for furthur usage..
Done..

```

0:01:20.626301

4.1.2 Build sample test data from the test data

In [41]:

```

start = datetime.now()

```



```

path = "sample_test_sparse_matrix_new.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk...")
    # just get it from the disk instead of computing it
    sample_test_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 5k users and 500 movies from available data
    sample_test_sparse_matrix = get_sample_sparse_matrix(test_sparse_matrix, no_users=5000, no_movies=500,
                                                         path = "sample_test_sparse_matrix_new.npz")

print(datetime.now() - start)

```

Original Matrix : (users, movies) -- (349312 17757)
Original Matrix : Ratings -- 20096102

Sampled Matrix : (users, movies) -- (5000 500)
Sampled Matrix : Ratings -- 7333
Saving it into disk for further usage..
Done..

0:00:20.031775

4.2 Finding Global Average of all movie ratings, Average rating per User, and Average rating per Movie (from sampled train)

In [42]:

```
sample_train_averages = dict()
```

4.2.1 Finding Global Average of all movie ratings

In [43]:

```

# get the global average of ratings in our train set.
global_average = sample_train_sparse_matrix.sum()/sample_train_sparse_matrix.count_nonzero()
sample_train_averages['global'] = global_average
sample_train_averages

```

Out[43]:

```
{'global': 3.5144714029443893}
```

4.2.2 Finding Average rating per User

In [44]:

```

sample_train_averages['user'] = get_average_ratings(sample_train_sparse_matrix, of_users=True)
print('\nAverage rating of user 2616751 :', sample_train_averages['user'][2616751])

```

Average rating of user 2616751 : 3.7777777777777777

4.2.3 Finding Average rating per Movie

In [45]:

```

sample_train_averages['movie'] = get_average_ratings(sample_train_sparse_matrix, of_users=False)
print('\nAverage rating of movie 46 :', sample_train_averages['movie'][46])

```

Average rating of movie 46 : 3.72992700729927

4.3 Featurizing data

In [46]:

```
print('\n No of ratings in Our Sampled train matrix is : {}'.format(sample_train_sparse_matrix.
count_nonzero()))
print('\n No of ratings in Our Sampled test  matrix is : {}'.format(sample_test_sparse_matrix.co
unt_nonzero()))
```

No of ratings in Our Sampled train matrix is : 199566

No of ratings in Our Sampled test matrix is : 7333

4.3.1 Featurizing data for regression problem

4.3.1.1 Featurizing train data

In [47]:

```
# get users, movies and ratings from our samples train sparse matrix
sample_train_users, sample_train_movies, sample_train_ratings =
sparse.find(sample_train_sparse_matrix)
```

In [48]:

```
#####
# It took me almost 10 hours to prepare this train dataset. #
#####
start = datetime.now()
if os.path.isfile('reg_train1.csv'):
    print("File already exists you don't have to prepare again..." )
else:
    print('preparing {} tuples for the dataset..{}\n'.format(len(sample_train_ratings)))
    with open('reg_train1.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_train_users, sample_train_movies,
sample_train_ratings):
            st = datetime.now()
            # print(user, movie)
            #----- Ratings of "movie" by similar users of "user" -----
            --
            # compute the similar Users of the "user"
            user_sim = cosine_similarity(sample_train_sparse_matrix[user],
sample_train_sparse_matrix).ravel()
            top_sim_users = user_sim.argsort()[::-1][1:] # we are ignoring 'The User' from its simi
lar users.
            # get the ratings of most similar users for this movie
            top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toarray().ravel()
            # we will make it's length "5" by adding movie averages to .
            top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_users_ratings.extend([sample_train_averages['movie'][movie]]*(5 -
len(top_sim_users_ratings)))
            # print(top_sim_users_ratings, end=" ")

            #----- Ratings by "user" to similar movies of "movie" -----
            ----
            # compute the similar movies of the "movie"
            movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie].T,
sample_train_sparse_matrix.T).ravel()
            top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignoring 'The User' from its si
milar users.
            # get the ratings of most similar movie rated by this user..
            top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toarray().ravel()
            # we will make it's length "5" by adding user averages to.
            top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_movies_ratings.extend([sample_train_averages['user']
```

```

[user]]*(5-len(top_sim_movies_ratings)))
# print(top_sim_movies_ratings, end=" : -- ")

#-----prepare the row to be stores in a file-----#
row = list()
row.append(user)
row.append(movie)
# Now add the other features to this data...
row.append(sample_train_averages['global']) # first feature
# next 5 features are similar users "movie" ratings
row.extend(top_sim_users_ratings)
# next 5 features are "user" ratings for similar_movies
row.extend(top_sim_movies_ratings)
# Avg_user rating
row.append(sample_train_averages['user'][user])
# Avg_movie rating
row.append(sample_train_averages['movie'][movie])

# finalley, The actual Rating of this user-movie pair...
row.append(rating)
count = count + 1

# add rows to the file opened..
reg_data_file.write(','.join(map(str, row)))
reg_data_file.write('\n')
if (count)%10000 == 0:
    # print(','.join(map(str, row)))
    print("Done for {} rows----- {}".format(count, datetime.now() - start))

print(datetime.now() - start)

```

preparing 199566 tuples for the dataset..

```

Done for 10000 rows----- 1:00:46.670212
Done for 20000 rows----- 1:58:49.779078
Done for 30000 rows----- 3:20:16.786290
Done for 40000 rows----- 4:31:07.128999
Done for 50000 rows----- 5:28:53.049387
Done for 60000 rows----- 6:28:06.063969
Done for 70000 rows----- 7:23:31.533590
Done for 80000 rows----- 8:18:50.768190
Done for 90000 rows----- 9:14:22.706638
Done for 100000 rows----- 10:14:51.522955
Done for 110000 rows----- 11:20:21.313238
Done for 120000 rows----- 12:28:51.408447
Done for 130000 rows----- 13:30:14.960645
Done for 140000 rows----- 14:26:27.009726
Done for 150000 rows----- 15:25:06.746113
Done for 160000 rows----- 16:23:38.578681
Done for 170000 rows----- 17:22:06.499123
Done for 180000 rows----- 18:25:17.616049
Done for 190000 rows----- 19:32:50.338636
20:39:58.488419

```

Reading from the file to make a Train_dataframe

In [49]:

```

reg_train = pd.read_csv('reg_train1.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2', 'sur3',
'sur4', 'sur5', 'smr1', 'smr2', 'smr3', 'smr4', 'smr5', 'UAvg', 'MAvg', 'rating'], header=None)
reg_train.head()

```

Out[49]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg	rating
0	286684	10	3.514471	3.0	2.0	4.0	3.0	3.0	3.0	3.0	3.0	3.0	3.0	2.830065	3.125	3
1	659540	10	3.514471	4.0	5.0	3.0	3.0	2.0	3.0	2.0	4.0	1.0	3.0	2.785714	3.125	3
2	1042361	10	3.514471	4.0	3.0	3.0	3.0	2.0	5.0	5.0	4.0	4.0	5.0	4.000000	3.125	5
3	1078512	10	3.514471	3.0	2.0	5.0	4.0	3.0	2.0	4.0	4.0	3.0	3.0	3.000000	3.125	2
4	1004517	10	3.514471	2.0	5.0	3.0	3.0	2.0	4.0	3.0	5.0	4.0	4.0	2.605652	3.125	4

4	1094517	10	3.514471	3.0	3.0	3.0	3.0	2.0	4.0	3.0	3.0	4.0	4.0	3.693692	3.125	4
	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg	rating

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
 - sur1, sur2, sur3, sur4, sur5 (top 5 similar users who rated that movie..)
- **Similar movies rated by this user:**
 - smr1, smr2, smr3, smr4, smr5 (top 5 similar movies rated by this movie..)
- **UAvg** : User's Average rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

4.3.1.2 Featurizing test data

In [50]:

```
# get users, movies and ratings from the Sampled Test
sample_test_users, sample_test_movies, sample_test_ratings = sparse.find(sample_test_sparse_matrix
)
```

In [51]:

```
sample_train_averages['global']
```

Out[51]:

```
3.5144714029443893
```

In [53]:

```
start = datetime.now()

if os.path.isfile('reg_test.csv'):
    print("It is already created...")
else:

    print('preparing {} tuples for the dataset..\n'.format(len(sample_test_ratings)))
    with open('reg_test.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_test_users, sample_test_movies,
sample_test_ratings):
            st = datetime.now()

            #----- Ratings of "movie" by similar users of "user" -----
            #print(user, movie)
            try:
                # compute the similar Users of the "user"
                user_sim = cosine_similarity(sample_train_sparse_matrix[user],
sample_train_sparse_matrix).ravel()
                top_sim_users = user_sim.argsort()[::-1][1:] # we are ignoring 'The User' from its
similar users.
                # get the ratings of most similar users for this movie
                top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toarray().ravel()
                # we will make it's length "5" by adding movie averages to .
                top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
                top_sim_users_ratings.extend([sample_train_averages['movie'][movie]]*(5 -
len(top_sim_users_ratings)))
                # print(top_sim_users_ratings, end="--")

            except (IndexError, KeyError):
                # It is a new User or new Movie or there are no ratings for given user for top sim:
lar movies...
```

```

##### Cold Start Problem #####
top_sim_users_ratings.extend([sample_train_averages['global']]*(5 -
len(top_sim_users_ratings)))
#print(top_sim_users_ratings)
except:
    print(user, movie)
    # we just want KeyErrors to be resolved. Not every Exception...
    raise

#----- Ratings by "user" to similar movies of "movie" -----
----

try:
    # compute the similar movies of the "movie"
    movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie].T,
sample_train_sparse_matrix.T).ravel()
    top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignoring 'The User' from it
s similar users.
    # get the ratings of most similar movie rated by this user..
    top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toarray().ravel()
    # we will make it's length "5" by adding user averages to.
    top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
    top_sim_movies_ratings.extend([sample_train_averages['user']
[user]]*(5-len(top_sim_movies_ratings)))
    #print(top_sim_movies_ratings)
except (IndexError, KeyError):
    #print(top_sim_movies_ratings, end=" : -- ")

top_sim_movies_ratings.extend([sample_train_averages['global']]*(5-len(top_sim_movies_ratings)))
#print(top_sim_movies_ratings)
except :
    raise

#-----prepare the row to be stores in a file-----#
row = list()
# add usser and movie name first
row.append(user)
row.append(movie)
row.append(sample_train_averages['global']) # first feature
#print(row)
# next 5 feaats are similar_users "movie" ratings
row.extend(top_sim_users_ratings)
#print(row)
# next 5 features are "user" ratings for similar_movies
row.extend(top_sim_movies_ratings)
#print(row)
# Avg_user rating
try:
    row.append(sample_train_averages['user'][user])
except KeyError:
    row.append(sample_train_averages['global'])
except:
    raise
#print(row)
# Avg_movie rating
try:
    row.append(sample_train_averages['movie'][movie])
except KeyError:
    row.append(sample_train_averages['global'])
except:
    raise
#print(row)
# finalley, The actual Rating of this user-movie pair...
row.append(rating)
#print(row)
count = count + 1

# add rows to the file opened..
reg_data_file.write(','.join(map(str, row)))
#print(','.join(map(str, row)))
reg_data_file.write('\n')
if (count)%1000 == 0:
    #print(','.join(map(str, row)))
    print("Done for {} rows----- {}".format(count, datetime.now() - start))
print("",datetime.now() - start)

```

preparing 7333 tuples for the dataset..

```
Done for 1000 rows----- 0:06:57.485448
Done for 2000 rows----- 0:14:06.727570
Done for 3000 rows----- 0:20:41.514729
Done for 4000 rows----- 0:26:35.054444
Done for 5000 rows----- 0:33:02.327582
Done for 6000 rows----- 0:39:29.617071
Done for 7000 rows----- 0:45:55.948803
0:48:03.377861
```

Reading from the file to make a test dataframe

In [60]:

```
reg_test_df = pd.read_csv('reg_test.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2', 'sur3',
'sur4', 'sur5',
                                                    'smr1', 'smr2', 'smr3', 'smr4', 'smr5',
'UAvg', 'MAvg', 'rating'], header=None)

reg_test_df.head(4)
```

Out[60]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	
0	808635	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471
1	941866	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471
2	1737912	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471
3	1849204	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
 - sur1, sur2, sur3, sur4, sur5 (top 5 similar users who rated that movie..)
- **Similar movies rated by this user:**
 - smr1, smr2, smr3, smr4, smr5 (top 5 similar movies rated by this movie..)
- **UAvg** : User AVerage rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

4.3.2 Transforming data for Surprise models

In [56]:

```
from surprise import Reader, Dataset
```

4.3.2.1 Transforming train data

- We can't give raw data (movie, user, rating) to train the model in Surprise library.
- They have a separate format for TRAIN and TEST data, which will be useful for training the models like SVD, KNNBaseLineOnly....etc., in Surprise.
- We can form the trainset from a file, or from a Pandas DataFrame.
http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py

In [57]:

```
# It is to specify how to read the dataframe.
# for our dataframe, we don't have to specify anything extra..
reader = Reader(rating_scale=(1,5))

# create the traindata from the dataframe...
train_data = Dataset.load_from_df(reg_train[['user', 'movie', 'rating']], reader)

# build the trainset from traindata.. It is of dataset format from surprise library..
trainset = train_data.build_full_trainset()
```

4.3.2.2 Transforming test data

- Testset is just a list of (user, movie, rating) tuples. (Order in the tuple is important)

In [61]:

```
testset = list(zip(reg_test_df.user.values, reg_test_df.movie.values, reg_test_df.rating.values))
testset[:3]
```

Out[61]:

```
[(808635, 71, 5), (941866, 71, 4), (1737912, 71, 3)]
```

4.4 Applying Machine Learning models

- Global dictionary that stores rmse and mape for all the models....
 - It stores the metrics in a dictionary of dictionaries

keys : model names(string)

value: dict(**key** : metric, **value** : value)

In [62]:

```
models_evaluation_train = dict()
models_evaluation_test = dict()

models_evaluation_train, models_evaluation_test
```

Out[62]:

```
(({}), {})
```

Utility functions for running regression models

In [63]:

```
# to get rmse and mape given actual and predicted ratings..
def get_error_metrics(y_true, y_pred):
    rmse = np.sqrt(np.mean([ (y_true[i] - y_pred[i])**2 for i in range(len(y_pred)) ]))
    mape = np.mean(np.abs( (y_true - y_pred)/y_true )) * 100
    return rmse, mape
```

```
#####
#####
def run_xgboost(algo, x_train, y_train, x_test, y_test, verbose=True):
    """
```

```

It will return train_results and test_results
"""

# dictionaries for storing train and test results
train_results = dict()
test_results = dict()

# fit the model
print('Training the model..')
start =datetime.now()
algo.fit(x_train, y_train, eval_metric = 'rmse')
print('Done. Time taken : {} \n'.format(datetime.now()-start))
print('Done \n')

# from the trained model, get the predictions....
print('Evaluating the model with TRAIN data...')
start =datetime.now()
y_train_pred = algo.predict(x_train)
# get the rmse and mape of train data...
rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)

# store the results in train_results dictionary..
train_results = {'rmse': rmse_train,
                 'mape' : mape_train,
                 'predictions' : y_train_pred}

#####
# get the test data predictions and compute rmse and mape
print('Evaluating Test data')
y_test_pred = algo.predict(x_test)
rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_test_pred)
# store them in our test results dictionary.
test_results = {'rmse': rmse_test,
               'mape' : mape_test,
               'predictions':y_test_pred}

if verbose:
    print('\nTEST DATA')
    print('-'*30)
    print('RMSE : ', rmse_test)
    print('MAPE : ', mape_test)

# return these train and test results...
return train_results, test_results

```

Utility functions for Surprise modes

In [64]:

```

# it is just to makesure that all of our algorithms should produce same results
# everytime they run...

my_seed = 15
random.seed(my_seed)
np.random.seed(my_seed)

#####
# get (actual_list , predicted_list) ratings given list
# of predictions (prediction is a class in Surprise).
#####
def get_ratings(predictions):
    actual = np.array([pred.r_ui for pred in predictions])
    pred = np.array([pred.est for pred in predictions])

    return actual, pred

#####
# get 'rmse' and 'mape' , given list of prediction objects
#####
def get_errors(predictions, print_them=False):

    actual, pred = get_ratings(predictions)

```



```

actual, pred = get_ratings(predictions)
rmse = np.sqrt(np.mean((pred - actual)**2))
mape = np.mean(np.abs(pred - actual)/actual)

return rmse, mape*100

#####
# It will return predicted ratings, rmse and mape of both train and test data #
#####
def run_surprise(algo, trainset, testset, verbose=True):
    """
        return train_dict, test_dict

        It returns two dictionaries, one for train and the other is for test
        Each of them have 3 key-value pairs, which specify 'rmse', 'mape', and 'predicted ratings'.
    """
    start = datetime.now()
    # dictionaries that stores metrics for train and test..
    train = dict()
    test = dict()

    # train the algorithm with the trainset
    st = datetime.now()
    print('Training the model...')
    algo.fit(trainset)
    print('Done. time taken : {} \n'.format(datetime.now()-st))

    # ----- Evaluating train data-----#
    st = datetime.now()
    print('Evaluating the model with train data..')
    # get the train predictions (list of prediction class inside Surprise)
    train_preds = algo.test(trainset.build_testset())
    # get predicted ratings from the train predictions..
    train_actual_ratings, train_pred_ratings = get_ratings(train_preds)
    # get 'rmse' and 'mape' from the train predictions.
    train_rmse, train_mape = get_errors(train_preds)
    print('time taken : {}'.format(datetime.now()-st))

    if verbose:
        print('-'*15)
        print('Train Data')
        print('-'*15)
        print("RMSE : {}\nMAPE : {}".format(train_rmse, train_mape))

    #store them in the train dictionary
    if verbose:
        print('adding train results in the dictionary..')
    train['rmse'] = train_rmse
    train['mape'] = train_mape
    train['predictions'] = train_pred_ratings

    #----- Evaluating Test data-----#
    st = datetime.now()
    print('\nEvaluating for test data...')
    # get the predictions( list of prediction classes) of test data
    test_preds = algo.test(testset)
    # get the predicted ratings from the list of predictions
    test_actual_ratings, test_pred_ratings = get_ratings(test_preds)
    # get error metrics from the predicted and actual ratings
    test_rmse, test_mape = get_errors(test_preds)
    print('time taken : {}'.format(datetime.now()-st))

    if verbose:
        print('-'*15)
        print('Test Data')
        print('-'*15)
        print("RMSE : {}\nMAPE : {}".format(test_rmse, test_mape))
    # store them in test dictionary
    if verbose:
        print('storing the test results in test dictionary...')
    test['rmse'] = test_rmse
    test['mape'] = test_mape
    test['predictions'] = test_pred_ratings

    print('\n'+ '-'*45)
    print('Total time taken to run this algorithm :', datetime.now() - start)

```

```
# return two dictionaries train and test
return train, test
```

4.4.1 XGBoost with initial 13 features

In [68]:

```
import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV
```

In [71]:

```
# prepare Train data
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# Prepare Test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

# initialize Our first XGBoost model...
param={
    'min_child_weight':[5,10,15],
    'max_depth':[5,10,20,50,100],
    'learning_rate':[0.001,0.05,0.5,1],
    'subsample':[0.8,0.5,0.2],
    'n_estimators':[2,10,25,50,100]
}

first_xgb = xgb.XGBRegressor()
clf = RandomizedSearchCV(first_xgb,param, cv=None,n_iter=10,verbose=2)
clf.fit(x_train,y_train)
print(clf.best_params_)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05, total=4.5s

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 4.4s remaining: 0.0s

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05, total=4.2s

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05, total=4.1s

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05, total=4.6s

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05

[CV] subsample=0.8, n_estimators=10, min_child_weight=5, max_depth=20, learning_rate=0.05, total=5.2s

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05, total= 17.3s

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05, total= 15.0s

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05, total= 15.8s

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05, total= 15.1s

[CV] subsample=0.2, n_estimators=50, min_child_weight=15, max_depth=20, learning_rate=0.05

[illegible]

```

0.9s
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5, total=
4.0s
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5, total=
3.8s
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5, total=
3.9s
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5, total=
3.7s
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=50, learning_rate=0.5, total=
3.9s
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5,
total= 9.7s
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5,
total= 10.1s
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5,
total= 9.6s
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5,
total= 9.8s
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5
[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=100, learning_rate=0.5,
total= 10.5s
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5, total=
47.8s
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5, total=
48.7s
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5, total=
48.2s
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5, total=
48.1s
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5
[CV] subsample=0.8, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.5, total=
47.6s

```

```
[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 17.1min finished
```

```
{'subsample': 0.2, 'n_estimators': 50, 'min_child_weight': 15, 'max_depth': 20, 'learning_rate': 0.05}
```

In [72]:

```

first_xgb = xgb.XGBRegressor(n_estimators=50,min_child_weight=15,max_depth=20,learning_rate=0.05)
train_results, test_results = run_xgboost(first_xgb, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['first_algo'] = train_results
models_evaluation_test['first_algo'] = test_results

xgb.plot_importance(first_xgb)
plt.show()

```

Training the model..

Done. Time taken : 0:00:27.298386

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

```

-----
-----

```

RMSE : 1.1418789815838117
MAPE : 32.38839863777117



4.4.2 Surprise BaselineModel

In [73]:

```
from surprise import BaselineOnly
```

Predicted_rating : (baseline prediction)

-

http://surprise.readthedocs.io/en/stable/basic_algorithms.html#surprise.prediction_algorithmr
seline_only.BaselineOnly

$$\hat{r}_{ui} = b_{ui} = \mu + b_u + b_i$$

- μ : Average of all trainings in training data.
- b_u : User bias
- b_i : Item bias (movie biases)

Optimization function (Least Squares Problem)

- http://surprise.readthedocs.io/en/stable/prediction_algorithms.html#baselines-estimates-c
onfiguration

$$\sum_{r_{ui} \in R_{train}} \left(r_{ui} - (\mu + b_u + b_i) \right)^2 + \lambda (b_u^2 + b_i^2)$$

[mimimize] $\{b_u, b_i\}$

In [75]:

```
# options are to specify..., how to compute those user and item biases  
bsl_options = {'method': 'sgd',  
               'learning_rate': .001}
```

```

    }
    bsl_algo = BaselineOnly(bsl_options=bsl_options)
    # run this algorithm.., It will return the train and test results..
    bsl_train_results, bsl_test_results = run_surprise(bsl_algo, trainset, testset, verbose=True)

    # Just store these error metrics in our models_evaluation datastructure
    models_evaluation_train['bsl_algo'] = bsl_train_results
    models_evaluation_test['bsl_algo'] = bsl_test_results

```

```

Training the model...
Estimating biases using sgd...
Done. time taken : 0:00:01.672886

Evaluating the model with train data..
time taken : 0:00:02.238498
-----
Train Data
-----
RMSE : 0.9553157605240995

MAPE : 30.8066420556683

adding train results in the dictionary..

Evaluating for test data...
time taken : 0:00:00.089946
-----
Test Data
-----
RMSE : 1.0720761061391755

MAPE : 34.19154502168587

storing the test results in test dictionary...

-----
Total time taken to run this algorithm : 0:00:04.005322

```

4.4.3 XGBoost with initial 13 features + Surprise Baseline predictor

Updating Train Data

In [76]:

```

# add our baseline_predicted value as our feature..
reg_train['bslpr'] = models_evaluation_train['bsl_algo']['predictions']
reg_train.head(2)

```

Out[76]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg	rating	bslpr
0	286684	10	3.514471	3.0	2.0	4.0	3.0	3.0	3.0	3.0	3.0	3.0	3.0	2.830065	3.125	3	2.856026
1	659540	10	3.514471	4.0	5.0	3.0	3.0	2.0	3.0	2.0	4.0	1.0	3.0	2.785714	3.125	3	3.204659

Updating Test Data

In [77]:

```

# add that baseline predicted ratings with Surprise to the test data as well
reg_test_df['bslpr'] = models_evaluation_test['bsl_algo']['predictions']

reg_test_df.head(2)

```

Out[77]:

	user user	movie movie	GAvg GAvg	sur1 sur1	sur2 sur2	sur3 sur3	sur4 sur4	sur5 sur5	smr1 smr1	smr2 smr2	smr3 smr3	smr4 smr4	smr5 smr5	U U
0	808635	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471
1	941866	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471

In [78]:

```
# prepare train data
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# Prepare Test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

# initialize Our first XGBoost model...
param={
    'min_child_weight':[5,10,15],
    'max_depth':[5,10,20,50,100],
    'learning_rate':[0.001,0.05,0.5,1],
    'subsample':[0.8,0.5,0.2],
    'n_estimators':[2,10,25,50,100]
}

xgb_bsl = xgb.XGBRegressor()
clf = RandomizedSearchCV(xgb_bsl,param, cv=None,n_iter=10,verbose=2)
clf.fit(x_train,y_train)
print(clf.best_params_)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits
[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1, total=7.8s

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 7.7s remaining: 0.0s

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1, total=7.5s

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1, total=7.8s

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1, total=7.4s

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1

[CV] subsample=0.8, n_estimators=25, min_child_weight=5, max_depth=10, learning_rate=1, total= 3.1s

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001, total= 2.6s

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001, total= 2.5s

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001, total= 2.5s

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001, total= 2.2s

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=10, learning_rate=0.001, total= 2.1s

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=5, learning_rate=1

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=5, learning_rate=1, total= 1.4s

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=5, learning_rate=1

[CV] subsample=0.2, n_estimators=10, min_child_weight=15, max_depth=5, learning_rate=1, total= 1.6s

[illegible]


```

[CV] subsample=0.2, n_estimators=2, min_child_weight=10, max_depth=50, learning_rate=1, total=
0.9s
[CV] subsample=0.2, n_estimators=2, min_child_weight=10, max_depth=50, learning_rate=1
[CV] subsample=0.2, n_estimators=2, min_child_weight=10, max_depth=50, learning_rate=1, total=
0.9s
[CV] subsample=0.2, n_estimators=2, min_child_weight=10, max_depth=50, learning_rate=1
[CV] subsample=0.2, n_estimators=2, min_child_weight=10, max_depth=50, learning_rate=1, total=
0.9s
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001, total=
4.8s
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001, total=
5.4s
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001, total=
5.1s
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001, total=
4.9s
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001
[CV] subsample=0.2, n_estimators=50, min_child_weight=5, max_depth=5, learning_rate=0.001, total=
5.2s
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5, total=
5.7s
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5, total=
5.5s
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5, total=
5.7s
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5, total=
5.9s
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5
[CV] subsample=0.2, n_estimators=10, min_child_weight=5, max_depth=50, learning_rate=0.5, total=
5.8s

```

```
[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 4.8min finished
```

```
{'subsample': 0.8, 'n_estimators': 100, 'min_child_weight': 10, 'max_depth': 5, 'learning_rate': 0.5}
```

In [80]:

```

xgb_bsl = xgb.XGBRegressor(subsample= 0.8, n_estimators= 100, min_child_weight= 10, max_depth= 5, l
earning_rate= 0.5
)

train_results, test_results = run_xgboost(xgb_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_bsl'] = train_results
models_evaluation_test['xgb_bsl'] = test_results

xgb.plot_importance(xgb_bsl)
plt.show()

```

Training the model..

Done. Time taken : 0:00:21.857580

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.149638266353395

MAPE : 32.99813931007494



4.4.4 Surprise KNNBaseline predictor

In [81]:

```
from surprise import KNNBaseline
```

- KNN BASELINE
 - http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithms.knns.KNNBaseline
- PEARSON_BASELINE SIMILARITY
 - http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_baseline
- SHRINKAGE
 - 2.2 Neighborhood Models in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>

• predicted Rating : (based on User-User similarity)

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{v \in N^k_i(u)} \text{sim}(u, v) \cdot (r_{vi} - b_{vi})}{\sum_{v \in N^k_i(u)} \text{sim}(u, v)}$$

- b_{ui} - Baseline prediction of (user, movie) rating
- $N^k_i(u)$ - Set of **K** similar users (neighbours) of **user (u)** who rated **movie(i)**
- $\text{sim}(u, v)$ - **Similarity** between users **u** and **v**
 - Generally, it will be cosine similarity or Pearson correlation coefficient.
 - But we use **shrunk Pearson-baseline correlation coefficient**, which is based on the pearsonBaseline similarity (we take base line predictions instead of mean rating of user/item)
- **Predicted rating** (based on Item Item similarity):

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in N^k_u(i)} \text{sim}(i, j) \cdot (r_{uj} - b_{uj})}{\sum_{j \in N^k_u(i)} \text{sim}(i, j)}$$
 - **Notations follows same as above (user user based predicted rating)**

4.4.4.1 Surprise KNNBaseline with user user similarities

In [82]:

```
# we specify , how to compute similarities and what to consider with sim_options to our algorithm
sim_options = {'user_based' : True,
               'name': 'pearson_baseline',
               'shrinkage': 100,
               'min_support': 2
              }

# we keep other parameters like regularization parameter and learning_rate as default values.
bsl_options = {'method': 'sgd'}

knn_bsl_u = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_options)
knn_bsl_u_train_results, knn_bsl_u_test_results = run_surprise(knn_bsl_u, trainset, testset,
verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_u'] = knn_bsl_u_train_results
models_evaluation_test['knn_bsl_u'] = knn_bsl_u_test_results
```

Training the model...
Estimating biases using sgd...
Computing the pearson_baseline similarity matrix...
Done computing similarity matrix.
Done. time taken : 0:02:28.024556

Evaluating the model with train data..
time taken : 0:05:47.669443

Train Data

RMSE : 0.3527396830995068

MAPE : 9.97148808270328

adding train results in the dictionary..

Evaluating for test data...
time taken : 0:00:00.151905

Test Data

RMSE : 1.0711072013527172

MAPE : 34.24892969621231

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:08:15.847903

4.4.4.2 Surprise KNNBaseline with movie movie similarities

In [83]:

```
# we specify , how to compute similarities and what to consider with sim_options to our algorithm

# 'user_based' : Fals => this considers the similarities of movies instead of users

sim_options = {'user_based' : False,
               'name': 'pearson_baseline',
               'shrinkage': 100,
               'min_support': 2
              }

# we keep other parameters like regularization parameter and learning_rate as default values.
bsl_options = {'method': 'sgd'}

knn_bsl_m = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_options)

knn_bsl_m_train_results, knn_bsl_m_test_results = run_surprise(knn_bsl_m, trainset, testset,
verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_m'] = knn_bsl_m_train_results
models_evaluation_test['knn_bsl_m'] = knn_bsl_m_test_results
```

```
models_evaluation_train['knn_bsl_m'] = knn_bsl_m_train_results
models_evaluation_test['knn_bsl_m'] = knn_bsl_m_test_results
```

```
Training the model...
Estimating biases using sgd...
Computing the pearson_baseline similarity matrix...
Done computing similarity matrix.
Done. time taken : 0:00:03.778465
```

```
Evaluating the model with train data..
time taken : 0:00:25.178241
```

```
-----
Train Data
```

```
-----
RMSE : 0.355275048093376
```

```
MAPE : 9.60667148221924
```

```
adding train results in the dictionary..
```

```
Evaluating for test data...
time taken : 0:00:00.103939
```

```
-----
Test Data
```

```
-----
RMSE : 1.0710588488031434
```

```
MAPE : 34.228374629059815
```

```
storing the test results in test dictionary...
```

```
-----
Total time taken to run this algorithm : 0:00:29.060645
```

4.4.5 XGBoost with initial 13 features + Surprise Baseline predictor + KNNBaseline predictor

- First we will run XGBoost with predictions from both KNN's (that uses User_User and Item_Item similarities along with our previous features.
- Then we will run XGBoost with just predictions from both knn models and predictions from our baseline model.

Preparing Train data

In [84]:

```
# add the predicted values from both knns to this dataframe
reg_train['knn_bsl_u'] = models_evaluation_train['knn_bsl_u']['predictions']
reg_train['knn_bsl_m'] = models_evaluation_train['knn_bsl_m']['predictions']

reg_train.head(2)
```

Out[84]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg	rating	bslpr	knn_bsl
0	286684	10	3.514471	3.0	2.0	4.0	3.0	3.0	3.0	3.0	3.0	3.0	3.0	2.830065	3.125	3	2.856026	3.0000
1	659540	10	3.514471	4.0	5.0	3.0	3.0	2.0	3.0	2.0	4.0	1.0	3.0	2.785714	3.125	3	3.204659	3.4062

Preparing Test data

In [85]:

```
reg_test_df['knn_bsl_u'] = models_evaluation_test['knn_bsl_u']['predictions']
reg_test_df['knn_bsl_m'] = models_evaluation_test['knn_bsl_m']['predictions']

reg_test_df.head(2)
```

Out[85]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	U
0	808635	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471
1	941866	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471

In [87]:

```
# prepare the train data....
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# prepare the train data....
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

# declare the model

param={
    'min_child_weight':[5,10,15],
    'max_depth':[5,10,20,50,100],
    'learning_rate':[0.001,0.05,0.5,1],
    'subsample':[0.8,0.5,0.2],
    'n_estimators':[2,10,25,50,100]
}

xgb_knn_bsl = xgb.XGBRegressor()
clf = RandomizedSearchCV(xgb_knn_bsl,param, cv=None,n_iter=10,verbose=2)
clf.fit(x_train,y_train)
print(clf.best_params_)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5, total= 45.2s

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 45.1s remaining: 0.0s

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5, total= 44.2s

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5, total= 46.4s

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5, total= 45.5s

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5

[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.5, total= 44.6s

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1, total= 16.9s

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1, total= 17.5s

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1, total= 16.3s

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1, total= 17.1s

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1

[CV] subsample=0.5, n_estimators=25, min_child_weight=10, max_depth=20, learning_rate=1, total= 17.4s

[CV] subsample=0.5, n_estimators=10, min_child_weight=15, max_depth=5, learning_rate=0.001

[CV] subsample=0.5, n_estimators=10, min_child_weight=15, max_depth=5, learning_rate=0.001,

[illegible]

```

[CV] subsample=0.5, n_estimators=10, min_child_weight=10, max_depth=50, learning_rate=1, total=
12.6s
[CV] subsample=0.5, n_estimators=10, min_child_weight=10, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=10, min_child_weight=10, max_depth=50, learning_rate=1, total=
11.6s
[CV] subsample=0.5, n_estimators=10, min_child_weight=10, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=10, min_child_weight=10, max_depth=50, learning_rate=1, total=
12.3s
[CV] subsample=0.5, n_estimators=10, min_child_weight=10, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=10, min_child_weight=10, max_depth=50, learning_rate=1, total=
11.2s
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total=
29.1s
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total=
28.3s
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total=
29.8s
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total=
29.2s
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1
[CV] subsample=0.5, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total=
28.4s
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05,
total= 43.4s
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05,
total= 44.5s
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05,
total= 44.4s
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05,
total= 42.6s
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05
[CV] subsample=0.2, n_estimators=100, min_child_weight=15, max_depth=20, learning_rate=0.05,
total= 44.0s

```

[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 15.5min finished

```
{'subsample': 0.2, 'n_estimators': 100, 'min_child_weight': 15, 'max_depth': 20, 'learning_rate':
0.05}
```

In [89]:

```

xgb_knn_bsl = xgb.XGBRegressor(subsample= 0.2, n_estimators= 100, min_child_weight= 15, max_depth=
20, learning_rate= 0.05
)
train_results, test_results = run_xgboost(xgb_knn_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_knn_bsl'] = train_results
models_evaluation_test['xgb_knn_bsl'] = test_results

xgb.plot_importance(xgb_knn_bsl)
plt.show()

```

Training the model..

Done. Time taken : 0:00:59.057824

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.081495350664273

MAPE : 33.771637642273575



4.4.6 Matrix Factorization Techniques

4.4.6.1 SVD Matrix Factorization User Movie interactions

In [90]:

```
from surprise import SVD
```

http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization.SVD

- Predicted Rating :

- $\hat{r}_{ui} = \mu + b_u + b_i + q_i^T p_u$
- q_i - Representation of item(movie) in latent factor space
- p_u - Representation of user in new latent factor space

- A BASIC MATRIX FACTORIZATION MODEL in [https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf)

- Optimization problem with user item interactions and regularization (to avoid overfitting)

$$\sum_{(r_{ui}) \in R_{\{train\}}} \left(r_{ui} - \hat{r}_{ui} \right)^2 +$$

$$\lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2)$$

In [91]:

```
# initialize the model
svd = SVD(n_factors=100, biased=True, random_state=15, verbose=True)
svd_train_results, svd_test_results = run_surprise(svd, trainset, testset, verbose=True)

# Just store these error metrics in our models evaluation datastructure
```



```
models_evaluation_train['svd'] = svd_train_results
models_evaluation_test['svd'] = svd_test_results
```

Training the model...

```
Processing epoch 0
Processing epoch 1
Processing epoch 2
Processing epoch 3
Processing epoch 4
Processing epoch 5
Processing epoch 6
Processing epoch 7
Processing epoch 8
Processing epoch 9
Processing epoch 10
Processing epoch 11
Processing epoch 12
Processing epoch 13
Processing epoch 14
Processing epoch 15
Processing epoch 16
Processing epoch 17
Processing epoch 18
Processing epoch 19
Done. time taken : 0:00:19.694456
```

Evaluating the model with train data..
time taken : 0:00:03.547662

Train Data

RMSE : 0.6635462679466172

MAPE : 20.59711721221903

adding train results in the dictionary..

Evaluating for test data..
time taken : 0:00:00.109936

Test Data

RMSE : 1.0716090944837267

MAPE : 34.167928194781176

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:00:23.357050

4.4.6.2 SVD Matrix Factorization with implicit feedback from user (user rated movies)

In [92]:

```
from surprise import SVDpp
```

- -----> 2.5 Implicit Feedback in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>

- Predicted Rating :

$$- \text{ } \hat{r}_{ui} = \mu + b_u + b_i + q_i^T \left(p_u + |I_u|^{-\frac{1}{2}} \sum_{j \in I_u} y_j \right)$$

- I_u --- the set of all items rated by user u
- y_j --- Our new set of item factors that capture implicit ratings.

- Optimization problem with user item interactions and regularization (to avoid overfitting)

$$- \sum_{\{ui\} \in R_{\{train\}}} (r_{\{ui\}} - \hat{r}_{\{ui\}})^2 +$$

$$\lambda(b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2 + \|y_j\|^2)$$

In [93]:

```
# initialize the model
svdpp = SVDpp(n_factors=50, random_state=15, verbose=True)
svdpp_train_results, svdpp_test_results = run_surprise(svdpp, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['svdpp'] = svdpp_train_results
models_evaluation_test['svdpp'] = svdpp_test_results
```

Training the model...

```
processing epoch 0
processing epoch 1
processing epoch 2
processing epoch 3
processing epoch 4
processing epoch 5
processing epoch 6
processing epoch 7
processing epoch 8
processing epoch 9
processing epoch 10
processing epoch 11
processing epoch 12
processing epoch 13
processing epoch 14
processing epoch 15
processing epoch 16
processing epoch 17
processing epoch 18
processing epoch 19
```

Done. time taken : 0:06:02.764824

Evaluating the model with train data..

time taken : 0:00:17.308800

Train Data

RMSE : 0.6128470138807365

MAPE : 18.48653191442907

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.106936

Test Data

RMSE : 1.0731604277442883

MAPE : 33.97520570315057

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:06:20.182556

4.4.7 XgBoost with 13 features + Surprise Baseline + Surprise KNNbaseline + MF Techniques

Preparing Train data

In [94]:

```
# add the predicted values from both knns to this dataframe
reg_train['svd'] = models_evaluation_train['svd']['predictions']
reg_train['svdpp'] = models_evaluation_train['svdpp']['predictions']

reg_train.head(2)
```

Out[94]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	...	smr4	smr5	UAvg	MAvg	rating	bslpr	knn_bsl_u
0	286684	10	3.514471	3.0	2.0	4.0	3.0	3.0	3.0	3.0	...	3.0	3.0	2.830065	3.125	3	2.856026	3.000000
1	659540	10	3.514471	4.0	5.0	3.0	3.0	2.0	3.0	2.0	...	1.0	3.0	2.785714	3.125	3	3.204659	3.406205

2 rows × 21 columns

Preparing Test data

In [95]:

```
reg_test_df['svd'] = models_evaluation_test['svd']['predictions']
reg_test_df['svdpp'] = models_evaluation_test['svdpp']['predictions']

reg_test_df.head(2)
```

Out[95]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	...	smr4	smr5	UAvg	
0	808635	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	...	3.514471	3.514471	3.514471	3
1	941866	71	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	3.514471	...	3.514471	3.514471	3.514471	3

2 rows × 21 columns

In [97]:

```
# prepare x_train and y_train
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# prepare test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

param={
    'min_child_weight':[5,10,15],
    'max_depth':[5,10,20,50,100],
    'learning_rate':[0.001,0.05,0.5,1],
    'subsample':[0.8,0.5,0.2],
    'n_estimators':[2,10,25,50,100]
}

xgb_final = xgb.XGBRegressor()
clf = RandomizedSearchCV(xgb_final,param, cv=None,n_iter=10)
clf.fit(x_train,y_train)
print(clf.best_params_)
```

```
{'subsample': 0.5, 'n_estimators': 50, 'min_child_weight': 10, 'max_depth': 10, 'learning_rate': 0.5}
```

In [98]:

In [98]:

```
xgb_final = xgb.XGBRegressor(subsample= 0.5, n_estimators= 50, min_child_weight= 10, max_depth= 10,
learning_rate= 0.5
)
train_results, test_results = run_xgboost(xgb_final, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_final'] = train_results
models_evaluation_test['xgb_final'] = test_results

xgb.plot_importance(xgb_final)
plt.show()
```

Training the model..

Done. Time taken : 0:00:24.027524

Done


Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

RMSE : 1.2154494709131842

MAPE : 31.183350174674885



4.4.8 XgBoost with Surprise Baseline + Surprise KNNbaseline + MF Techniques

In [99]:

```
# prepare train data
x_train = reg_train[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_train = reg_train['rating']

# test data
x_test = reg_test_df[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_test = reg_test_df['rating']

param={
    'min_child_weight':[5,10,15],
    'max_depth':[5,10,20,50,100],
    'learning_rate':[0.001,0.05,0.5,1],
    'subsample':[0.5,0.8,1.0],
    'n_estimators':[50,100,200]
```

```
'subsample':[0.8,0.5,0.2],
'n_estimators':[2,10,25,50,100]
}
```

```
xgb_all_models = xgb.XGBRegressor()
clf = RandomizedSearchCV(xgb_all_models,param, cv=None,n_iter=10,verbose=2)
clf.fit(x_train,y_train)
print(clf.best_params_)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05, total= 24.4s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 24.3s remaining: 0.0s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05, total= 20.8s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05, total= 19.6s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05, total= 20.4s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=20, learning_rate=0.05, total= 19.4s

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total= 8.8s

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total= 8.6s

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total= 8.0s

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total= 8.6s

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=25, min_child_weight=15, max_depth=50, learning_rate=1, total= 9.1s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1, total= 42.1s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1, total= 41.4s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1, total= 42.3s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1, total= 40.6s

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1

[CV] subsample=0.2, n_estimators=100, min_child_weight=5, max_depth=50, learning_rate=1, total= 41.7s

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05, total= 0.3s

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05, total= 0.4s

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05, total= 0.4s

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05

[CV] subsample=0.5, n_estimators=2, min_child_weight=5, max_depth=10, learning_rate=0.05, total= 0.4s

[illegible]

```

0.4s
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001,
total= 2.6s
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001,
total= 2.3s
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001,
total= 2.2s
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001,
total= 2.0s
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001
[CV] subsample=0.8, n_estimators=25, min_child_weight=15, max_depth=5, learning_rate=0.001,
total= 2.1s

```

```
[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 9.8min finished
```

```
{'subsample': 0.5, 'n_estimators': 2, 'min_child_weight': 10, 'max_depth': 5, 'learning_rate': 1}
```

In [100]:

```

xgb_all_models = xgb.XGBRegressor(subsample= 0.5, n_estimators= 2, min_child_weight= 10, max_depth=
5, learning_rate= 1
)
train_results, test_results = run_xgboost(xgb_all_models, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_all_models'] = train_results
models_evaluation_test['xgb_all_models'] = test_results

xgb.plot_importance(xgb_all_models)
plt.show()

```

Training the model..

Done. Time taken : 0:00:00.496692

Done


Evaluating the model with TRAIN data...


Evaluating Test data

TEST DATA

RMSE : 1.0773084995761828

MAPE : 34.68929096917531





4.5 Comparison between all models

In [101]:

```
# Saving our TEST_RESULTS into a dataframe so that you don't have to run it again
pd.DataFrame(models_evaluation_test).to_csv('small_sample_results.csv')
models = pd.read_csv('small_sample_results.csv', index_col=0)
models.loc['rmse'].sort_values()
```

Out[101]:

```
knn_bsl_m          1.0710588488031434
knn_bsl_u          1.0711072013527172
svd                 1.0716090944837267
bsl_algo           1.0720761061391755
svdpp              1.0731604277442883
xgb_all_models     1.0773084995761828
xgb_knn_bsl        1.081495350664273
first_algo         1.1418789815838117
xgb_bsl            1.149638266353395
xgb_final          1.2154494709131842
Name: rmse, dtype: object
```

5. Assignment

1. Instead of using 10K users and 1K movies to train the above models, use 25K users and 3K movies (or more) to train all of the above models. Report the RMSE and MAPE on the test data using larger amount of data and provide a comparison between various models as shown above.

NOTE: Please be patient as some of the code snippets make take many hours to complete execution.

2. Tune hyperparameters of all the Xgboost models above to improve the RMSE.

In []:

```
%%javascript
// Converts integer to roman numeral
// https://github.com/kmahelona/ipython_notebook_goodies
// https://kmahelona.github.io/ipython_notebook_goodies/ipython_notebook_toc.js
function romanize(num) {
    var lookup = {M:1000,CM:900,D:500,CD:400,C:100,XC:90,L:50,XL:40,X:10,IX:9,V:5,IV:4,I:1},
    roman = '',
    i;
    for ( i in lookup ) {
        while ( num >= lookup[i] ) {
            roman += i;
            num -= lookup[i];
        }
    }
    return roman;
}

// Builds a <ul> Table of Contents from all <headers> in DOM
function createTOC() {
    var toc = "";
    var level = 0;
    var levels = {}
    $('#toc').html('');

    $(':header').each(function(i) {
        if (this.id=='tocheading') {return;}
    })
}
```



```

    var titleText = this.innerHTML;
    var openLevel = this.tagName[1];

    if (levels[openLevel]){
        levels[openLevel] += 1;
    } else{
        levels[openLevel] = 1;
    }

    if (openLevel > level) {
        toc += (new Array(openLevel - level + 1)).join('<ul class="toc">');
    } else if (openLevel < level) {
        toc += (new Array(level - openLevel + 1)).join("</ul>");
        for (i=level;i>openLevel;i--){levels[i]=0;}
    }

    level = parseInt(openLevel);

    if (this.id==''){this.id = this.innerHTML.replace(/ /g,"-")}
    var anchor = this.id;

    toc += '<li><a style="text-decoration:none", href="#" + encodeURIComponent(anchor) + ">' + titleText + '</a></li>';

});

    if (level) {
        toc += (new Array(level + 1)).join("</ul>");
    }

    $('#toc').append(toc);

};

// Executes the createToc function
setTimeout(function(){createTOC();},100);

// Rebuild to TOC every minute
setInterval(function(){createTOC();},60000);

```